

Q2

```
In [1]: ### Load the dataset (which is saved as a pickle file)
import numpy as np
import matplotlib.pyplot as plt
import pickle


with open('dataset.pkl', 'rb') as f: # Python 3: open(..., 'rb')
    x_train, y_train, x_test, y_test = pickle.load(f)


# Note that each data point is a row
print('x_train has shape:', np.shape(x_train))
print('x_test has shape:', np.shape(x_test))

### Interactive scatter plot of dataset
%matplotlib notebook
from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure(figsize=(8, 8))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x_train[:,0], x_train[:,1], x_train[:,2], c=y_train, cmap='
```

x_train has shape: (2000, 3)
x_test has shape: (1018, 3)





Out[1]: <mpl_toolkits.mplot3d.art3d.Path3DCollection at 0x7ffbfc5e6bd0>

In [2]: *### From an earlier activity*
compute the mean of the three classes, return a column vector

```
# complete the code below
# hint 1 -- x_train[y_train==1,:] for example will extract only the el
# hint 2 -- np.mean(blah, axis=0) will take the mean of each row
# hint 3 -- reshape your vector so that it's a column vector
mu_0 = np.mean(x_train[y_train==0,:], axis=0).reshape((-1,1))
mu_1 = np.mean(x_train[y_train==1,:], axis=0).reshape((-1,1))
mu_2 = np.mean(x_train[y_train==2,:], axis=0).reshape((-1,1))

### compute covariance of each class
### np.cov() expects each column to be a datapoint
cov_0 = np.cov(x_train[y_train==0,:].T)
cov_1 = np.cov(x_train[y_train==1,:].T)
cov_2 = np.cov(x_train[y_train==2,:].T)
```

In [3]: *### complete the code below to compute the log-likelihood ratio under*
def log_likelihood(_x, _mu, _cov):
 ## _x and _mu should be column vectors, and _cov should be an n \t
 assert np.shape(_x) == np.shape(_mu)
 _log_likelihood = -1*np.linalg.slogdet(_cov)[1] - (_x-_mu).T@np.li
return _log_likelihood[0,0]

```
In [4]: from sklearn.metrics import classification_report
### predict the class of the vectors in the test set
y_hat = []
for i, x in enumerate(x_test):
    x_column_vector = np.reshape(x, (-1,1))
    ll0 = log_likelihood(x_column_vector, mu_0, cov_0)
    ll1 = log_likelihood(x_column_vector, mu_1, cov_1)
    ll2 = log_likelihood(x_column_vector, mu_2, cov_2)
    y_hat.append(np.argmax([ll0, ll1, ll2]))

### compute the accuracy and print a classification report
print(classification_report(y_test, y_hat))
```


	precision	recall	f1-score	support
0	0.98	0.92	0.95	341
1	0.90	0.96	0.93	336
2	0.97	0.95	0.96	341
accuracy			0.94	1018
macro avg	0.95	0.95	0.95	1018
weighted avg	0.95	0.94	0.95	1018

```
In [5]: ### create data points from three classes, and plot for comparison
x_0 = np.random.multivariate_normal(mu_0.squeeze(), cov_0, 1000)
x_1 = np.random.multivariate_normal(mu_1.squeeze(), cov_1, 1000)
x_2 = np.random.multivariate_normal(mu_2.squeeze(), cov_2, 1000)
print(np.shape(x_0))

# %matplotlib notebook #uncomment this line to make plot interactive
from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure(figsize=(8, 8))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x_0[:,0], x_0[:,1], x_0[:,2], c='r', cmap='jet')
ax.scatter(x_1[:,0], x_1[:,1], x_1[:,2], c='b', cmap='jet')
ax.scatter(x_2[:,0], x_2[:,1], x_2[:,2], c='g', cmap='jet')
```

(1000, 3)



Out[5]: <mpl_toolkits.mplot3d.art3d.Path3DCollection at 0x7ffbfe817d50>


a

```
In [77]: from sklearn import mixture  
  
### Start here.  
### Write code below to implement Gaussian mixture and EM
```

```
In [78]: gmm = mixture.GaussianMixture(n_components=3, covariance_type='full')  
gmm.fit(x_train)  
y_EM_train = gmm.fit_predict(x_train)
```

b

```
In [79]: fig = plt.figure(figsize=(8, 8))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x_train[:,0], x_train[:,1], x_train[:,2] , c=y_EM_train, cm
```



Out[79]: <mpl_toolkits.mplot3d.art3d.Path3DCollection at 0x7ffbfe874c50>

C

In [80]: # ML

```
print("params computed from ML")
print("means: ", np.hstack((mu_0, mu_1, mu_2)).T)
print("covariances: ", cov_0)
print(cov_1)
print(cov_2)
```

```
params computed from ML
means: [[-1713.81050084  944.31365051  281.16236007]
 [-1552.93107583  -89.19746319  477.87368097]
 [-1295.33434171  -205.02751531  -567.1897333 ]]
covariances: [[[ 2.00183376e+05 -6.52227338e+04  1.36975829e+02]
 [-6.52227338e+04  1.37326374e+05  4.47994782e+04]
 [ 1.36975829e+02  4.47994782e+04  6.68313191e+04]]
[[189841.36785878 -12635.95079816 -38014.69527155]
 [-12635.95079816  71312.25210536  34677.35464327]
 [-38014.69527155  34677.35464327  107049.81922793]]
[[102911.65495427   8789.1205339   35885.2850861 ]
 [ 8789.1205339   139937.25639306 -43562.24330135]
 [ 35885.2850861   -43562.24330135  98891.78375075]]]
```

In [81]: #EM

```
print("means: ", gmm.means_)
print('covariances: ', gmm.covariances_)
```

```
means: [[-1288.0444108  -173.0877544  -653.01294118]
 [-1699.55379161   966.96794663   266.44530813]
 [-1537.80757098  -136.74988108  415.8095794 ]]
covariances: [[[ 92926.18576186   7392.86839683  40796.13896518]
 [ 7392.86839683  111629.0675653  -20887.38717239]
 [ 40796.13896518  -20887.38717239  54858.73908749]]
[[204055.92443845  -72202.038599  -4986.79477082]
 [-72202.038599   112880.81277548  45961.71237651]
 [-4986.79477082   45961.71237651  78372.65897091]]
[[191665.19009974  -23188.2726937  -38221.56943699]
 [-23188.2726937   97330.10131853  53259.91724865]
 [-38221.56943699   53259.91724865  123375.14331804]]]
```

Results from 2 algorithms deviate from each other

d

```
In [97]: y_EM_test = gmm.predict(x_test)

# re-assign labels
y_EM_test = np.where(y_EM_test==2, -1, y_EM_test)
y_EM_test = np.where(y_EM_test==0, -2, y_EM_test)
y_EM_test = np.where(y_EM_test==1, 0, y_EM_test)
y_EM_test = np.abs(y_EM_test)

err_count = np.sum(y_EM_test != y_test)
print("error rate = ", err_count / len(y_test))
```

error rate = 0.08546168958742632

e

```
In [109]: from sklearn import cluster
km = cluster.KMeans(n_clusters=3)
km.fit(x_train)
y_KM_train = km.fit_predict(x_train)
```

```
In [116]: y_KM_test = km.predict(x_test)

# re-assign labels
y_KM_test = np.where(y_KM_test==0, -2, y_KM_test)
# y_KM_test = np.where(y_KM_test==0, -2, y_KM_test)
y_KM_test = np.where(y_KM_test==2, 0, y_KM_test)
y_KM_test = np.abs(y_KM_test)

err_count = np.sum(y_KM_test != y_test)
print("error rate = ", err_count / len(y_test))
```

error rate = 0.06581532416502947

error rate slightly higher than EM

f

Try cluster/mixture counts from 1 to, say 10 for MNIST, then pick the iteration with min error rate

Q3

a

In []: