

```
In [1]: import numpy as np
import tensorflow as tf
from sklearn.linear_model import LogisticRegression
import matplotlib.pyplot as plt
from matplotlib import pyplot as plt

(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
print(np.shape(x_train))

def vectorize(_image):
    return np.reshape(_image, (-1, 1))

vec_x_train = np.squeeze(np.array([vectorize(m) for m in x_train]))
vec_x_test = np.squeeze(np.array([vectorize(m) for m in x_test]))

# generate an instance of the logistic regression class model with multiple classes
model = LogisticRegression(solver='saga', tol=0.01, multi_class='multinomial')

(60000, 28, 28)
```

```
In [2]: # use model.fit(some arguments) to fit the model to the data
model.fit(vec_x_train, y_train)
```

```
Out[2]: LogisticRegression(multi_class='multinomial', solver='saga', tol=0.01)
```

```
In [4]: from sklearn.metrics import classification_report
```

```
### compute the accuracy and print a classification report
# use model.predict(some arguments) to predict the class of the test data
y_head_test = model.predict(vec_x_test)

# use classification_report(some arguments) to print accuracy
print(classification_report(y_test, y_head_test))
```

	precision	recall	f1-score	support
0	0.95	0.98	0.97	980
1	0.96	0.98	0.97	1135
2	0.93	0.90	0.91	1032
3	0.90	0.91	0.91	1010
4	0.93	0.93	0.93	982
5	0.91	0.86	0.88	892
6	0.95	0.95	0.95	958
7	0.93	0.92	0.93	1028
8	0.88	0.88	0.88	974
9	0.91	0.91	0.91	1009
accuracy			0.93	10000
macro avg	0.92	0.92	0.92	10000
weighted avg	0.93	0.93	0.93	10000

```
In [5]: from sklearn.preprocessing import OneHotEncoder
```

```
# convert the 10 classes to one hot encoding
one_hot = OneHotEncoder()
Y_train = one_hot.fit_transform(y_train.reshape(-1,1)).toarray()
Y_test = one_hot.fit_transform(y_test.reshape(-1,1)).toarray()
print(np.shape(Y_train))
```

```
(60000, 10)
```

```
In [30]: import keras
from keras.models import Sequential
from keras.layers import Dense

model = Sequential()
model.add(Dense(10, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=[accuracy])

# use model.fit(some arguments) to fit the model to the data, specify
model.fit(vec_x_train, Y_train, batch_size = 5, epochs = 10)
```

```
Epoch 1/10
60000/60000 [=====] - 15s 255us/step - loss: 14.3045 - accuracy: 0.8502
Epoch 2/10
60000/60000 [=====] - 15s 252us/step - loss: 13.3905 - accuracy: 0.8782
Epoch 3/10
60000/60000 [=====] - 15s 256us/step - loss: 13.4908 - accuracy: 0.8827
Epoch 4/10
60000/60000 [=====] - 16s 267us/step - loss: 13.2772 - accuracy: 0.8834
Epoch 5/10
60000/60000 [=====] - 17s 287us/step - loss: 13.3067 - accuracy: 0.8868
Epoch 6/10
60000/60000 [=====] - 18s 308us/step - loss: 13.2110 - accuracy: 0.8868
Epoch 7/10
60000/60000 [=====] - 19s 317us/step - loss: 13.3363 - accuracy: 0.8880
Epoch 8/10
60000/60000 [=====] - 20s 339us/step - loss: 13.3057 - accuracy: 0.8898
Epoch 9/10
60000/60000 [=====] - 17s 277us/step - loss: 13.3966 - accuracy: 0.8881
Epoch 10/10
60000/60000 [=====] - 15s 258us/step - loss: 13.3410 - accuracy: 0.8889
```

Out[30]: <keras.callbacks.callbacks.History at 0x7fbe5c951350>

```
In [32]: from sklearn.metrics import classification_report
```

```
### compute the accuracy and print a classification report
# use model.predict(some arguments) to predict the class of the test data
# use classification_report(some arguments) to print accuracy
Y_head_test = model.predict_classes(vec_x_test)
print(classification_report(y_test, Y_head_test))
```

	precision	recall	f1-score	support
0	0.95	0.97	0.96	980
1	0.98	0.90	0.94	1135
2	0.78	0.93	0.85	1032
3	0.75	0.93	0.83	1010
4	0.95	0.82	0.88	982
5	0.97	0.63	0.76	892
6	0.93	0.94	0.93	958
7	0.94	0.89	0.91	1028
8	0.85	0.81	0.83	974
9	0.80	0.92	0.86	1009
accuracy			0.88	10000
macro avg	0.89	0.87	0.87	10000
weighted avg	0.89	0.88	0.88	10000

Logistic regression model is more accurate and runs faster.